Introduction to ML
Exercise 4
Due Date: December 28th 22:00, 2021

Yosi Shrem, Yael Segal, Roni Chernyak and Yossi Keshet

December 14, 2021

Guidelines

1. You are allowed to work in pairs. If you choose this option, only one of the students should submit the exercise files.

2. You are allowed to use numpy and scipy packages and PyTorch framework.

3. Technical questions about this exercise should be asked at the course' piazza or during the TIRGUL.

4. Personal issues regarding the deadline should be directed to **Yael Segal**.

5. In order to submit your solution please submit the following files:

   (a) `details.txt` - A text file with your full name (in the first line) and ID (in the second line). See attached details file for pairs submission.

   (b) `ex_4.py` - A python 3.6+ file that contains your main function (attach ANY additional files needed for your code to run).

   (c) `ex_4_report.pdf` - A pdf file in which you describe your model and parameters.

   (d) `test_y` - your model's predictions on the given test set (see instructions below).

Follow the instructions and submit all files needed for your code to run.

Good Luck!

Ex4

In this exercise you will implement, train and evaluate your neural network using PyTorch package.

Installation   First, you will need to install PyTorch package. Installation instructions were uploaded to the Piazza. Please follow them and discuss issues there.

Data - FashionMNIST   This dataset contains 10 different categories of clothing. Each image is 28 pixels in height and 28 pixels in width, for a total of 784 pixels in total. Each pixel has a single pixel-value associated with it, indicating the lightness or darkness of that pixel. This pixel-value is an integer between 0 and 255.

Labels.   The possible labels are:

   0. T-shirt/top

   1. Trouser

   2. Pullover

   3. Dress

   4. Coat

   5. Sandal

   6. Shirt

   7. Sneaker

   8. Bag

   9. Ankle boot
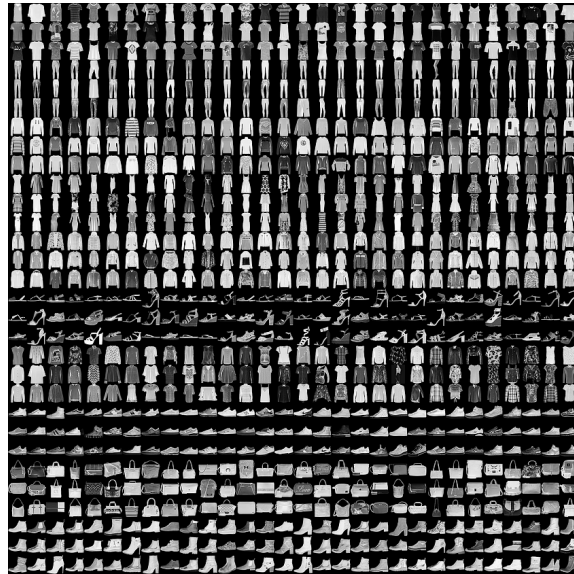
Instructions

In this exercise you will implement fully connected neural networks via Py-Torch. You will need to implement several settings and report the effect of each setting in terms of loss and accuracy. You should explore the following:

1. Model A - Neural Network with two hidden layers, the first layer should have a size of 100 and the second layer should have a size of 50, both should be followed by ReLU activation function. Train this model with SGD optimizer.

2. Model B - Neural Network with two hidden layers, the first layer should have a size of 100 and the second layer should have a size of 50, both should be followed by ReLU activation function, train this model with ADAM optimizer.

3. Model C - Dropout – add dropout layers to model B. You should place the dropout on the output of the hidden layers.

4. Model D - Batch Normalization - add Batch Normalization layers to model B. Where should you place the Batch Normalization layer? before the activation functions or after? try both cases and report the results in the report file.

5. Model E - Neural Network with five hidden layers:[128,64,10,10,10] using **ReLU** .

6. Model F - Neural Network with five hidden layers:[128,64,10,10,10] using **Sigmoid**.

In all these experiments you should use log_softmax as the output of the network and nll_loss function (see code example in recitation 7 slides).

Training

You should train your models using FashionMNIST dataset. You should train your models for 10 epochs each. You can use the code example we provide you in recitation 7 or in the PyTorch examples repository on GitHub.

You should split the training set to train and validation (80:20). You can load the provided data files using `train_x = numpy.loadtxt("train_x")` (repeat this for any data file you wish to load). Suggestion: save a small portion of the dataset and use it for debugging. When you are done, load the entire dataset to train your model.

Finally, you should use your best model (doesn't have to be one of the previous models) to generate predictions for the examples in `test_x` and write them into a file named `test_y`, similarly to the previous exercise. Your predictions file should contain 5000 rows exactly. Note: `Do not shuffle the test file`.

The run command to your program should be:

```
$ python ex4.py <train_x_path> <train_y_path>  <test_x_path> <output_log_name>
```

```
For example:
$ python ex4.py train_x train_y test_x test_y
```

Note: `Your program should run only your best model training`.

Evaluation - Report

Your report file, `ex_4_report.pdf`, should include the following for EACH model:

1. Plot the average **loss** per epoch for the validation and training set in a single image.

2. Plot the average **accuracy** per epoch for the validation and training set in a single image.

3. Test set accuracy (original FashionMNIST test set).

4. Hyper parameters.

Good Luck!